

# CIS520 - Project 4: Performance Analysis

Members: Benjamin Cooper, Bethany Weddle, Mir Fahim Anwar

Date: May 14, 2020

## Pthread Observation and Description:

In the observation of runtimes, it appears that the more cores the better, although having more threads does not make for faster runtime all the time. An accident test was run with 444 threads with 4 cores, but the runtime was not far from the result of 4 threads with 4 cores. **Bold = (Best Column), Bold & Underline = (Best Row), Purple = Shared**

Pthreads was the simplest to understand and implement. The pthread\_create and join made it really simple to start the parallelism and join all the threads to work together. We accounted for synchronization by having a start and end point for the threads that is calculated by the MAX\_THREADS, and the id of the thread that is passed by pthread\_create method to the FindSums method. The results do not outprint in order of line numbers, but the results are correct and consistent.

## Pthread Runtime (ms):

| Pthread     | Cores: 1,<br>Nodes: 1   | Cores: 2,<br>Nodes: 1   | Cores: 4,<br>Nodes: 1  | Cores: 8,<br>Nodes: 1   | Cores: 16,<br>Nodes: 1  | Cores: 16,<br>Nodes: 2  |
|-------------|-------------------------|-------------------------|------------------------|-------------------------|-------------------------|-------------------------|
| Threads: 1  | 16088.589               | 15040.830               | 15454.845              | 14206.231               | <b><u>13646.415</u></b> | <b><u>13723.885</u></b> |
| Threads: 2  | 16169.847               | 15102.817               | 14776.748              | <b><u>13387.157</u></b> | 13809.272               | 14003.610               |
| Threads: 4  | 17047.464               | 15399.775               | 15590.479              | 14258.153               | <b><u>13894.574</u></b> | 14431.968               |
| Threads: 8  | 16571.219               | 15621.84                | 15643.283              | <b><u>13316.115</u></b> | 13740.383               | 15037.787               |
| Threads: 16 | <b><u>16059.787</u></b> | <b><u>14566.352</u></b> | 13899.864              | 14581.026               | 13916.952               | <b><u>13735.689</u></b> |
| Threads: 32 | 16938.594               | 14734.625               | <b><u>13481.28</u></b> | <b><u>13437.882</u></b> | 13945.842               | 14181.025               |


## First 50 Lines Output:

line 0-1: -32923  
line 1-2: 144845  
line 2-3: -54368  
line 3-4: 71061  
line 4-5: -2101  
line 5-6: 2187  
line 25-26: -150417  
line 26-27: 141415  
line 27-28: -3493  
line 28-29: -127309  
line 29-30: -12523  
line 6-7: 3137  
line 7-8: -6344  
line 8-9: 4350  
line 9-10: -4558  
line 10-11: -3039  
line 30-31: 145174


line 31-32: -28979  
line 32-33: -104102  
line 33-34: 117276  
line 34-35: 15487  
line 35-36: -136449  
line 36-37: 144032  
line 37-38: 1779  
line 11-12: -135433  
line 12-13: 135721  
line 13-14: -154007  
line 38-39: -156143  
line 39-40: 135656  
line 40-41: -139426  
line 41-42: -5371  
line 42-43: -2978  
line 43-44: 72240  
line 44-45: 63548

line 45-46: -101234  
line 46-47: 87363  
line 47-48: 47957  
line 48-49: -6798  
line 49-50: 3632  
line 14-15: 164651  
line 15-16: -156641  
line 16-17: 4546  
line 17-18: -14809  
line 18-19: 26528  
line 19-20: -6322  
line 20-21: 138054  
line 21-22: -35903  
line 22-23: 14261  
line 23-24: -130131  
line 24-25: 159709

Y-Axis Represents Milliseconds



Here, we clearly see that the parallel use of multiple CPUs consistently resulted in better runtimes compared to single-core execution. However, this effect does not scale for larger numbers of cores. For instance, our results show that there is no consistent difference between 8 and 16 core executions.



This graph shows there is a clear downward trend in runtime for increasing numbers of cores, up to about 8.

### OpenMP Observation and Description:

In runtimes again appear to not be overly affected by the number of threads, but with the increase of cores, the runtimes become smaller. It also appears that OpenMP is slower than Pthreads in its execution of the program. **Bold = (Best Column)**, **Bold & Underline = (Best Row)**, Purple = Shared

OpenMP's code went through multiple revisions of trying to read the file in parallel as well as finding the sum differences. This did not work out because the threads kept interrupting each other to do the work, and the data would be inconsistent. So, we changed it to read the file into memory, and then traverse through an array of sums to show the differences. By OpenMP made parallel simple to start with the #pragma omp parallel, although it took many trial and error before we found the solution to synchronizing the threads.


### OpenMP Runtime (ms):

| OpenMP      | Cores: 1,<br>Nodes: 1 | Cores: 2,<br>Nodes: 1 | Cores: 4,<br>Nodes: 1 | Cores: 8,<br>Nodes: 1 | Cores: 16,<br>Nodes: 1  | Cores: 16,<br>Nodes: 2  |
|-------------|-----------------------|-----------------------|-----------------------|-----------------------|-------------------------|-------------------------|
| Threads: 1  | 18309.102             | <b>16383.792</b>      | 16677.015             | 15299.809             | <u>14244.468</u>        | 15444.925               |
| Threads: 2  | 17240.266             | 16451.943             | 15853.849             | <b>14347.395</b>      | <b><u>14279.383</u></b> | 15593.810               |
| Threads: 4  | 17155.521             | 16494.000             | 16599.727             | 14395.587             | 14271.495               | <b><u>13702.234</u></b> |
| Threads: 8  | <b>17097.161</b>      | 16805.256             | <b>15350.794</b>      | 14372.312             | 14293.426               | <u>13662.298</u>        |
| Threads: 16 | 17727.337             | 16423.814             | 16056.780             | 15461.847             | <b><u>14398.244</u></b> | 15600.227               |
| Threads: 32 | 17437.464             | 16720.983             | 16697.874             | 15437.794             | 14479.559               | <b><u>13694.473</u></b> |


### First 50 Lines Output:

|                     |                     |                     |
|---------------------|---------------------|---------------------|
| line 0-1: 32923     | line 17-18: 14809   | line 34-35: -15487  |
| line 1-2: -144845   | line 18-19: -26528  | line 35-36: 136449  |
| line 2-3: 54368     | line 19-20: 6322    | line 36-37: -144032 |
| line 3-4: -71061    | line 20-21: -138054 | line 37-38: -1779   |
| line 4-5: 2101      | line 21-22: 35903   | line 38-39: 156143  |
| line 5-6: -2187     | line 22-23: -14261  | line 39-40: -135656 |
| line 6-7: -3137     | line 23-24: 130131  | line 40-41: 139426  |
| line 7-8: 6344      | line 24-25: -159709 | line 41-42: 5371    |
| line 8-9: -4350     | line 25-26: 150417  | line 42-43: 2978    |
| line 9-10: 4558     | line 26-27: -141415 | line 43-44: -72240  |
| line 10-11: 3039    | line 27-28: 3493    | line 44-45: -63548  |
| line 11-12: 135433  | line 28-29: 127309  | line 45-46: 101234  |
| line 12-13: -135721 | line 29-30: 12523   | line 46-47: -87363  |
| line 13-14: 154007  | line 30-31: -145174 | line 47-48: -47957  |
| line 14-15: -164651 | line 31-32: 28979   | line 48-49: 6798    |
| line 15-16: 156641  | line 32-33: 104102  | line 49-50: -3632   |
| line 16-17: -4546   | line 33-34: -117276 |                     |

Y-Axis Represents Milliseconds



This graph shows the same general trends as the one for Pthreads.



This graph shows that the running times do trend downward as the cores increase. Although using multiple nodes sometimes results in worse performance.

### MPI Observation and Description:

MPI was the most frustrating to implement because running it on the head node appeared to display the correct differences for the lines, although once submitting jobs on beocat, the output would contain the data displaying multiple times. This was causing the jobs to take exceptionally longer times than we expected. It turns out we were using “mpirun” in the slurm launch script, which caused slurm to run  $n^2$  copies of the process instead of  $n$ . Correcting this small syntax error, allowed us to run our MPI program with varying core/node configurations.


MPI does not explicitly use threads in the same way as pthread and openMP do. Parallelism is implemented by running copies of a single binary as separate processes, usually on separate CPU cores. So for the MPI jobs, we specified the cores and the number of nodes among which those cores are distributed to get the following results.

### MPI Runtime (ms):

| MPI | Cores: 1<br>Node: 1 | Cores: 2<br>Node: 1 | Cores: 4<br>Node: 1 | Cores: 8<br>Node: 1 | Cores: 16<br>Node: 1,2 | Cores: 32<br>Node: 2 | Cores: 64<br>Node: 4 |
|-----|---------------------|---------------------|---------------------|---------------------|------------------------|----------------------|----------------------|
|     | 39048.937           | 38613.447           | 36648.339           | 37586.933           | 38702.484              | 43460.442            | 45356.267            |
|     |                     |                     |                     |                     | 38883.399              |                      |                      |

### First 50 Lines Output:

|                |                |                |
|----------------|----------------|----------------|
| 0-1: -32923    | 17-18: -14809  | 34-35: 15487   |
| 1-2: 144845    | 18-19: 26528   | 35-36: -136449 |
| 2-3: -54368    | 19-20: -6322   | 36-37: 144032  |
| 3-4: 71061     | 20-21: 138054  | 37-38: 1779    |
| 4-5: -2101     | 21-22: -35903  | 38-39: -156143 |
| 5-6: 2187      | 22-23: 14261   | 39-40: 135656  |
| 6-7: 3137      | 23-24: -130131 | 40-41: -139426 |
| 7-8: -6344     | 24-25: 159709  | 41-42: -5371   |
| 8-9: 4350      | 25-26: -150417 | 42-43: -2978   |
| 9-10: -4558    | 26-27: 141415  | 43-44: 72240   |
| 10-11: -3039   | 27-28: -3493   | 44-45: 63548   |
| 11-12: -135433 | 28-29: -127309 | 45-46: -101234 |
| 12-13: 135721  | 29-30: -12523  | 46-47: 87363   |
| 13-14: -154007 | 30-31: 145174  | 47-48: 47957   |
| 14-15: 164651  | 31-32: -28979  | 48-49: -6798   |
| 15-16: -156641 | 32-33: -104102 | 49-50: 363     |
| 16-17: 4546    | 33-34: 117276  |                |



This graph shows how the pthread program is consistently faster on average than openMP for all thread and cores combinations.



From this graph, we can see that our implementation of MPI failed to successfully use parallelism to speed up processing. The runtimes remained fairly consistent across multiple core/node configurations. We believe the slight upward trend in runtime at higher core counts is due to the overhead incurred by coordinating multiple processes.

## Pthread Code:

```
#include <stdlib.h>
#include <stdio.h>
#include <stdint.h>
#include <sys/time.h>
#include <pthread.h>
#define MAX_THREADS 32
#define BUFFSIZE 1000000
#define FILENAME "/homes/dan/625/wiki_dump.txt"

int sums[BUFFSIZE];
int count = 0;

void *FindSums(void *arg)
{
    int j;
    int id = (uintptr_t)arg;
    int start = id * (BUFFSIZE / MAX_THREADS);
    int end = start + (BUFFSIZE / MAX_THREADS);

    // j needs to start where thread left off from
    for(j = start; j < end; j++){
        printf("tid-%d line %d-%d: %d\n", pthread_self(), j,j+1,(sums[j+1]-sums[j]));
    }

    pthread_exit((void*) arg);
}

int main()
{
    struct timeval start, end;
    double elapsedTime;
    int numSlots, myVersion = 2; // pthreads = 1, openmp = 2, mpi = 3

    FILE *fp;
    char c = 0; // To store a character read from file
    int sum = 0;
    int i = 0;

    int t, rc;
    pthread_t threads[MAX_THREADS];
    pthread_attr_t attr;
    void *status;
    /* Initialize and set thread detached attribute */
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);

    gettimeofday(&start, NULL);
    // Open the file
    fp = fopen(FILENAME, "r");
```

```

// Check if file exists
if (fp == NULL)
{
    printf("Could not open file"); // breaks here
    return 0;
}

// Extract characters from file and store in character c
for(i=0; i <= BUFSIZE; i++) {
    c = getc(fp);
    while(c != '\n') {
        if(c == EOF) {
            break;
        }
        sum += (int)c;
        c = getc(fp);
    }
    sums[i] = sum;
    count++;
    sum = 0;
}

for (t = 0; t < MAX_THREADS; t++) {
    rc = pthread_create(&threads[t], &attr, FindSums, (void *) (uintptr_t)t);
    if (rc)
    {
        printf("ERROR; return code from pthread_create() is %d\n", rc);
        exit(-1);
    }
}

pthread_attr_destroy(&attr);
for(t=0; t< MAX_THREADS; t++)
{
    rc = pthread_join(threads[t], &status);
    if (rc)
    {
        printf("ERROR; return code from pthread_join() is %d\n", rc);
        exit(-1);
    }
}

gettimeofday(&end, NULL);

elapsedTime = (end.tv_sec - start.tv_sec) * 1000.0; //sec to ms
elapsedTime += (end.tv_usec - start.tv_usec) / 1000.0; // us to ms
printf("DATA, %d, %s, %f, %d\n", myVersion, getenv("NSLOTS"), elapsedTime, MAX_THREADS );

fclose(fp);
return 0;
}

```



## OpenMP Code:

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/time.h>
#include <omp.h>
#define MAX_THREADS 32
#define NUMLINES 1000000
#define FILENAME "/homes/dan/625/wiki_dump.txt"

int main() {
    struct timeval start, end;
    double elapsedTime;
    int numSlots, myVersion = 1; // omp = 1, pthread = 2, mpi = 3

    FILE * fp;
    int count = 0; // tracks total number of lines read
    char c = 0; // stores the char read from file
    int sum = 0; // the sum of a line's chars
    int * sums = malloc(NUMLINES * sizeof(int)); // a buffer to hold line sums
    int i = 0;
    int j = 0;

    omp_set_num_threads(MAX_THREADS);

    // Open the file
    fp = fopen(FILENAME, "r");

    // Check if file exists
    if (fp == NULL) {
        printf("Failed to open file");
        return -1;
    }
    printf("File opened successfully\n");

    gettimeofday(&start, NULL);
    // loop which reads characters from the file, stopping when EOF is reached or buffer is full
    while(c = getc(fp), c != EOF && i < NUMLINES) {
        if (c == '\n') {
            sums[i] = sum; // save line sum to sums buffer
            i++;
            sum = 0; // reset value for next line
        } else {
            sum += (int)c; // add each char to line sum
        }
    }
    count += i; // count tracks total number of lines across all loops

#pragma omp parallel
    {
        #pragma omp for ordered
        for (j = 0; j < count-1; j++) {
            #pragma omp ordered
            printf("(thread %d) line %d-%d: %d\n", omp_get_thread_num(), j, j + 1, (sums[j + 1] - sums[j]));
        }
    }

    gettimeofday(&end, NULL);

    elapsedTime = (end.tv_sec - start.tv_sec) * 1000.0; //sec to ms
    elapsedTime += (end.tv_usec - start.tv_usec) / 1000.0; // us to ms
    printf("DATA, %d, %s, %f, %d\n", myVersion, getenv("NSLOTS"), elapsedTime, MAX_THREADS);

    fclose(fp);
    return 0;
}
```

## MPI Code:

```
#include <stdlib.h>
#include <stdio.h>
#include <mpi.h>
#include <sys/time.h>

#define BUFFSIZE 1000000
#define FILENAME "/homes/dan/625/wiki_dump.txt"

int num_proc;
int sums[BUFFSIZE];

void ReadFile()
{
    FILE *fp;
    int i = 0;
    int sum = 0;
    char c = 0;

    // Open the file
    fp = fopen(FILENAME, "r");

    // Check if file exists
    if (fp == NULL)
    {
        printf("Could not open file"); // breaks here
    }

    // Extract characters from file and store in character c
    for(i=0; i <= BUFFSIZE; i++) {
        c = getc(fp);
        while(c != '\n') {
            if(c == EOF) {
                break;
            }
            sum += (int)c;
            c = getc(fp);
        }
        sums[i] = sum;
        sum = 0;
    }

    fclose(fp);
}

void FindSums(int *rank)
{
    int j;
    int id = *rank;
    int start = id * (BUFFSIZE / num_proc);
    int end = start + (BUFFSIZE / num_proc);
```

```

    for(j = start; j < end; j++)
    {
        printf("%d-%d: %d\n", j, j+1, sums[j+1]-sums[j]);
    }
}

int main(int argc, char* argv[])
{
    struct timeval start, end;
    double elapsedTime;
    int numSlots, myVersion = 3;

    int m, rc;
    int numtasks, rank;
    MPI_Status Status;

    gettimeofday(&start, NULL);
    rc = MPI_Init(&argc, &argv);
    if(rc != MPI_SUCCESS) {
        printf ("Error starting MPI program. Terminating.\n");
        MPI_Abort(MPI_COMM_WORLD, rc);
    }

    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if(rank == 0){
        ReadFile();
    }

    num_proc = numtasks;

    MPI_Bcast(sums, BUFFSIZE, MPI_INT, 0, MPI_COMM_WORLD);
    FindSums(&rank);
    MPI_Barrier(MPI_COMM_WORLD);

    MPI_Finalize();
    if(rank == 0) {
        gettimeofday(&end, NULL);
        elapsedTime = (end.tv_sec - start.tv_sec) * 1000.0; //sec to ms
        elapsedTime += (end.tv_usec - start.tv_usec) / 1000.0; // us to ms
        printf("DATA, %d, %s, %f, %d\n", myVersion, getenv("SLURM_CPUS_ON_NODE"),
elapsedTime, num_proc);
    }

    //MPI_Finalize();
    return 0;
}

```

## Links:

[https://github.com/ChristianJHughes/pintos-project4/blob/master/sample\\_code/proj4\\_code\\_beta/MPI\\_C\\_SAMPLE.c](https://github.com/ChristianJHughes/pintos-project4/blob/master/sample_code/proj4_code_beta/MPI_C_SAMPLE.c)  
<https://www.openmp.org/wp-content/uploads/OpenMP3.0-SummarySpec.pdf>  
[https://github.com/rpwilliams/CIS520\\_Proj4/blob/master/src/3way-mpi/lcs-mpi.c](https://github.com/rpwilliams/CIS520_Proj4/blob/master/src/3way-mpi/lcs-mpi.c)  
<https://support.beocat.ksu.edu/BeocatDocs/index.php>  
<https://linux.die.net/man/>  
<https://stackoverflow.com/questions/51414697/order-of-threads-executing-in-a-multiple-thread-program>  
<https://stackoverflow.com/questions/42887315/avoiding-race-condition-using-int-to-void-casting>  
<https://stackoverflow.com/questions/1845482/what-is-uintptr-t-data-type>  
<https://stackoverflow.com/questions/12490347/how-to-force-openmp-to-run-iterations-in-specific-order>  
[https://princetonuniversity.github.io/PUbootcamp/sessions/parallel-programming/Intro\\_PP\\_bootcamp\\_2018.pdf](https://princetonuniversity.github.io/PUbootcamp/sessions/parallel-programming/Intro_PP_bootcamp_2018.pdf)  
<https://github.com/wesleykendall/mpitutorial/blob/gh-pages/tutorials/mpi-send-and-receive/code/ring.c>  
<https://mpitutorial.com/tutorials/>  
<https://stackoverflow.com/questions/8841069/critical-section-in-mpi>